

The Compiler, the Runtime and other interesting "beasts" from the Clojure codebase

EuroClojure 2014



Renzo Borgatti - [@reborg](#)

MailOnline

Welcome to this talk about Clojure. The talk is not about the language from the end-user perspective but about the Clojure implementation itself. The implementation we are looking at is of course the reference RCI (Rich's Clojure Implementation). There is a lot of complexity involved in the implementation of such a successful language and this talk is my attempt to understand some of most interesting corner of the Clojure sources.

An amazing growth

Mar 2006: First commit

Oct 2006: 30k LOC 7 months old

Oct 2007: Clojure announced!

Oct 2008: Clojure invited at Lisp50

May 2009: 1.0 + book!

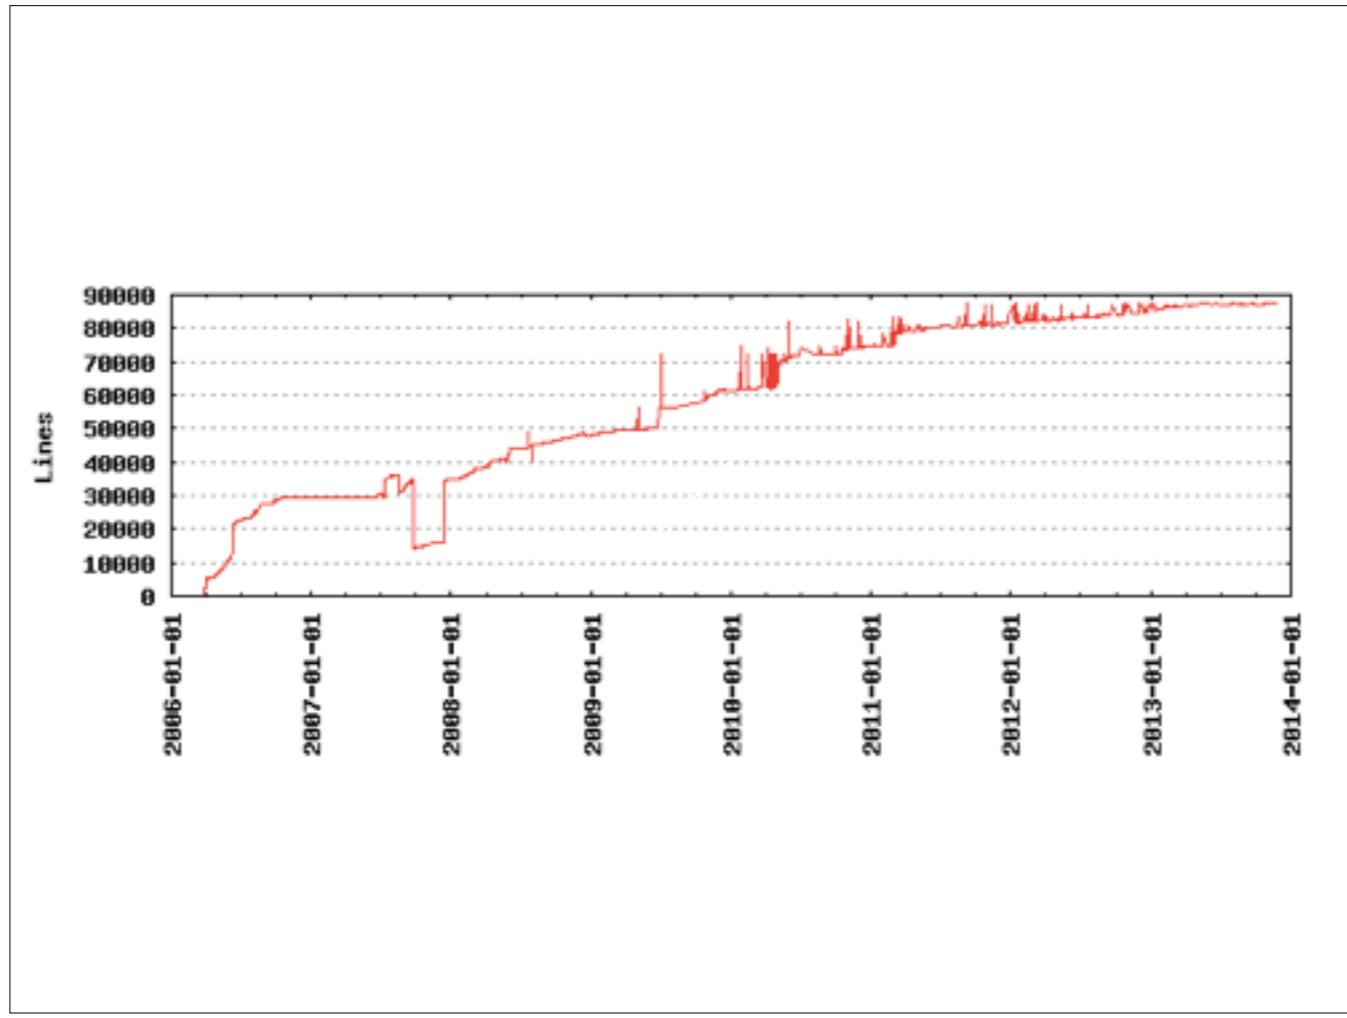
This is a quick overview of the important initial milestones of Clojure. One thing overall: the growth ratio is impressive: 30k LOC after only 7 months of work.

If we ever decide to have a Clojure retreat day, that must definitely be October the 10th.

Info:

<http://clojure.org/space/stats/overview>

<http://clojure-log.n01se.net>



Sorry for the brightness bomb.

Initial Code Milestones

Apr '06: Lisp2Java sources

May '06: boot.clj appears

May '06: STM first cut

Jun '06: First persistent data structure

Sep '06: Java2Java sources

Aug '07: Java2Bytecode started

Right after: almost all the rest, Refs, LockingTx

Clojure starts its existence as a Lisp DSL. The first Lisp implementation requires LispWorks to run and is generating compilable Java.

lib.lisp is the first incarnation of core.clj and contains the first implementation of the Clojure standard library.

All important language aspects we are used today are present since the very beginning. Work on persistence data structures, STM, Java interop proceeds almost in parallel.

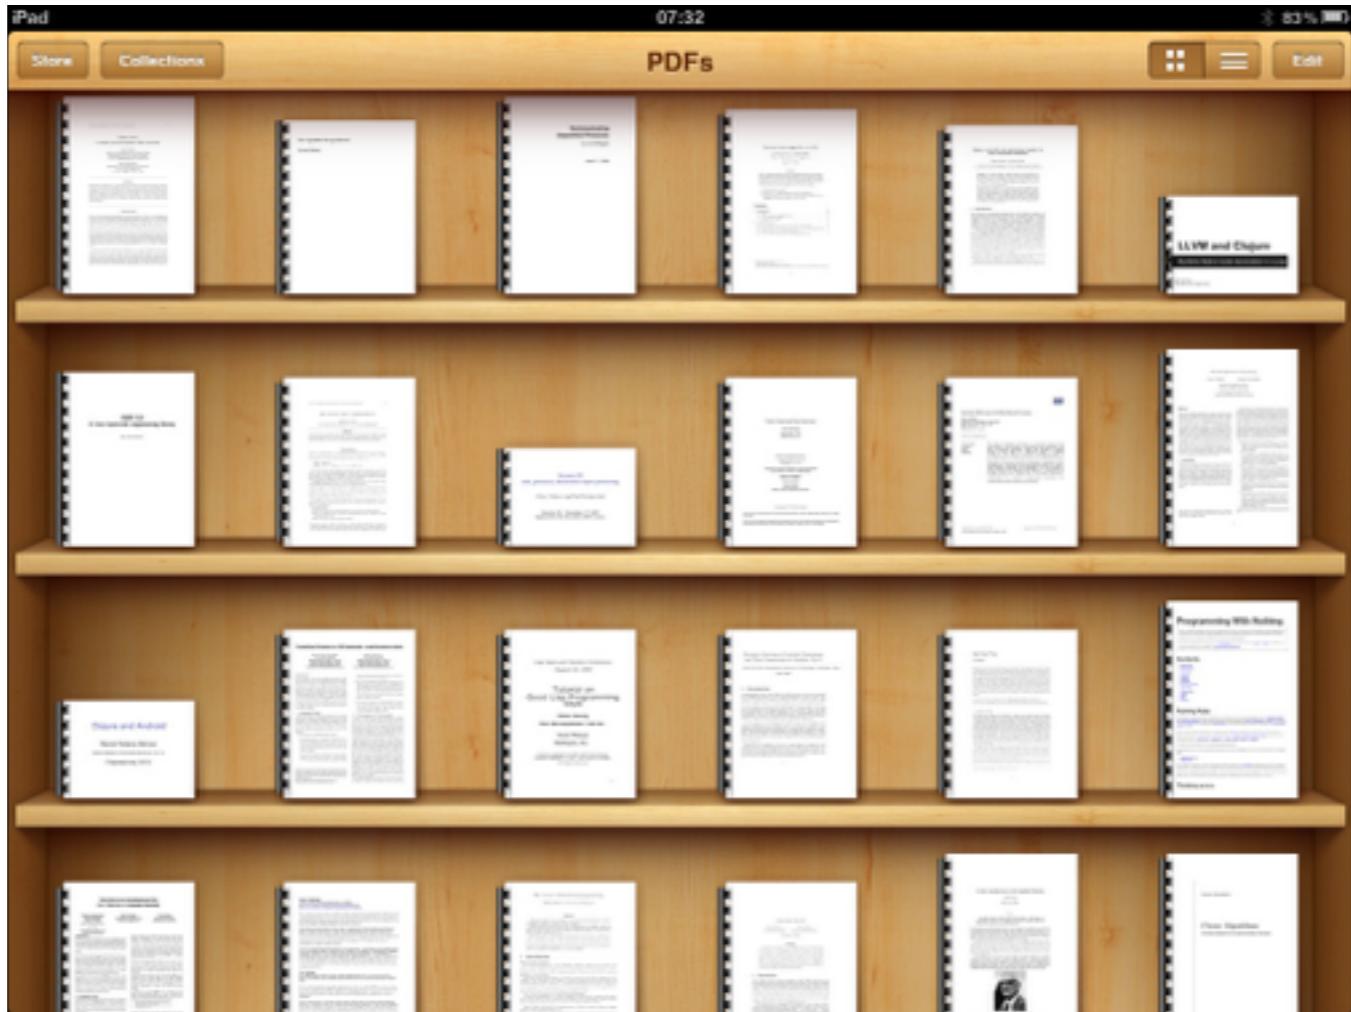
The introduction of the LispReader Java class is the first translation of parts of the Lisp code into Java. The next big part of it being the Compiler.

The Compiler first cut is mainly a translation of the Lisp code and is still generating Java sources.

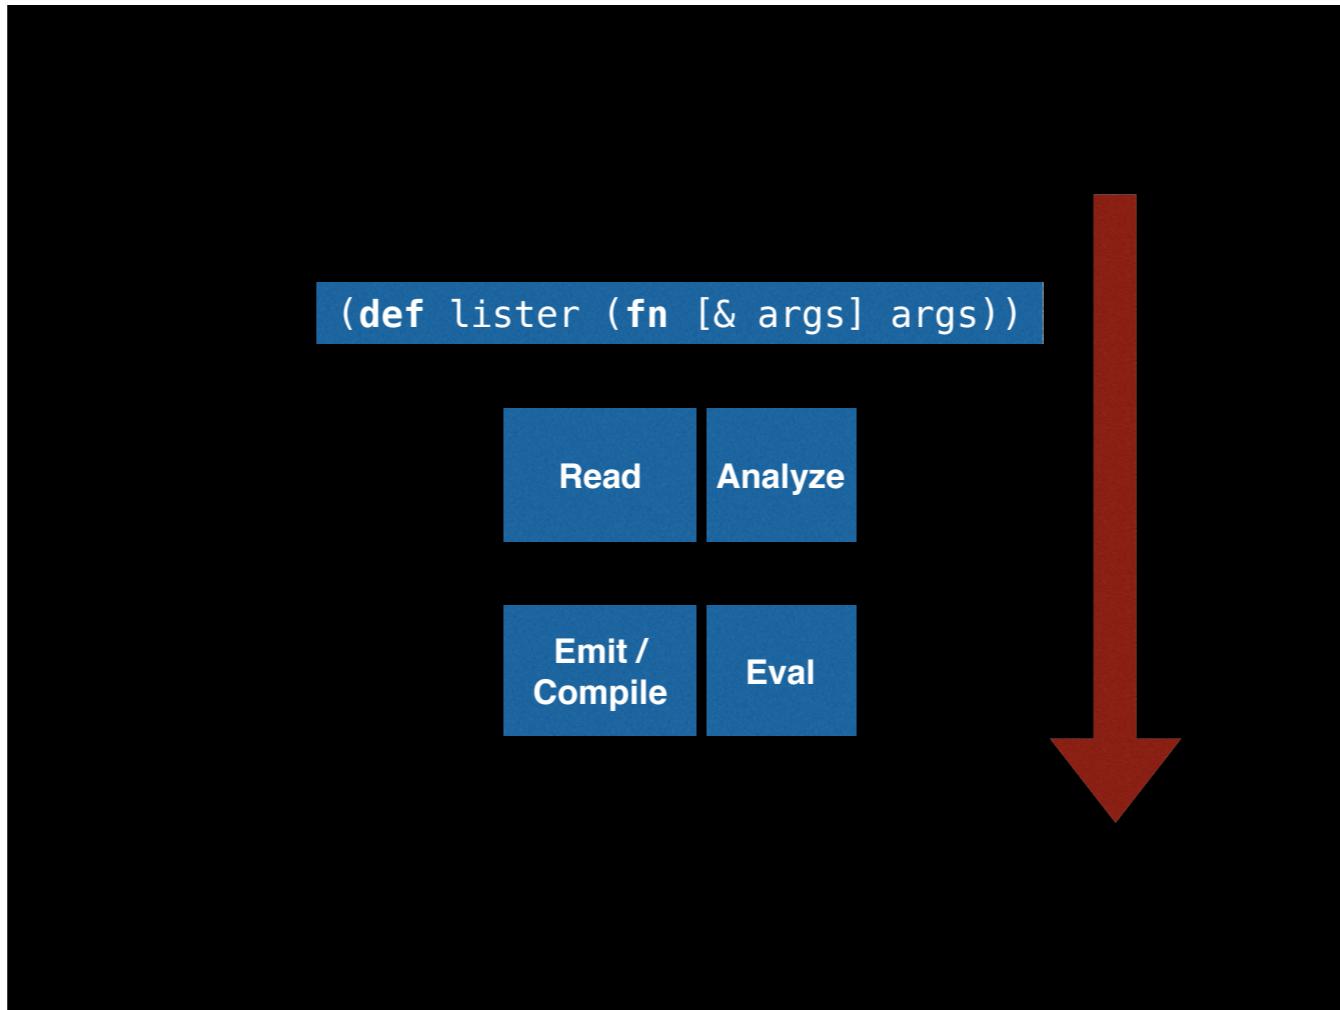
Curiosity: Clojure at some point got an ANTLR grammar experiment, if you ever interested in one (although maybe not up to date) you need to dig the commits.

The introduction of the BytecodeCompiler is architecturally the last big design change. With the BytecodeCompiler Clojure can finally have a REPL and be truly dynamic.

Of course there is much more since 2007 to nowadays and things got terribly more complicated, but the same compiler architecture is still there with Clojure happily generating tons of byte code fun.



iBook on my iPad, 4 months into studying the Clojure sources. Go refresh your CS degree! Now! Data structures and algorithms are pervasive. Deep JVM concurrency knowledge. Lisp internals, concepts and heritage. Big OO hierarchies and deep FP



30,000 ft view. Clojure can be thought very broadly as a big Lisp DSL to drive byte code generation. We can quickly see this in action by dumping the class that is generated for a function. In order to achieve that result, Clojure needs to parse the original form, understand what instructions are contained in it and emit the necessary bytecode. In the next section we are going to explore the main objects involved during processing of Clojure s-expressions. S-expr means symbolic expression, a form of notation to express recursive lists of “symbols” introduced with Lisp. From the point of view of parsing textual s-expressions Clojure is very similar to Lisp.

Reader

```
(def lister (fn [& args] args))
```

```
(  
  def  
  lister  
  (  
    fn  
    [& args]  
    args))
```

```
PersistentList(  
  Symbol("def"),  
  Symbol("lister"),  
  PersistentList(  
    Symbol("fn"),  
    PersistentVector("&", Symbol("args")),  
    Symbol("args")))
```

By the way, ever wondered why there is a LispReader in Clojure that reads Clojure forms? Isn't this a ClojureReader? Some of the early Clojure history is showing here. Let's take a simple example. We feed Clojure (it can be at the REPL or loading a script with that line in it) with a definition for a function of a variable number of arguments that returns those arguments as a list. The list function indeed. The LispReader decomposes text (which is always a list in Clojure) into instances of internal classes, like PersistentList of Symbols. The resulting outer PersistentList is usually called a "form". The nice thing about recursive list is that they can be handled recursively, from parsing to analysing and then to the emission phase.

- Transforms Clojure text into symbolic lists
- Cons, PersistentVector, IteratorSeq and so on
- Uses a push back reader to read/unread
- Perform part of the syntax validation

Analyzer

```
(def lister (fn [& args] args))
```

```
PersistentList(  
  Sym("def"),  
  Sym("lister"),  
  PersistentList(  
    Sym("fn"),  
    PersistentVector(  
      "&",  
      Sym("args"),  
      Sym("args"))))
```

```
DefExpr(  
  Var("lister"),  
  FnExpr(  
    Sym("lister"),  
    PersistentList(  
      FnMethod(  
        LocalBinding(Sym("args")),  
        BodyExpr(  
          PersistentVector(  
            LocalBindingExpr("args"))))))))
```

- Transform LispReader forms into Expr
- Compiler.analyze() & friends
- Expr inner classes implementing IParser

There is not really an Analyzer class in Clojure. The compiler contains 4 analyzer functions. The rest of the analysis is done by implementors of the IParser interface, which are all static inner classes of Expr interface implementors. Expr classes are for example: DefExpr, FnExpr, MethodExpr, BodyExpr, LocalBindingExpr and so on.

Let's continue on the previous expansion. Based on the (first PersistentList) Symbol (def, fn, and so on) a specific Parser class is selected from a corresponding Expr class. The parser instance knows how to build expression objects out of the (rest PersistentList) or stop on a leaf without invoking "analyze" any further. The result is a recursive hierarchy of Expr classes that know how to emit/eval them selves.

Emission

Mainly bytecode generation for Expr

Prerequisite for evaluation

emit() method in Expr interface

Notable exception: called over FnExpr analysis

Emitting has to do with generation of bytecode. The emit() method takes a context, the expression to transform and an ASM::GeneratorAdapter instance. For example DefExpr.emit() will coordinate bytecode generation for its meta info and it's initialisation form. Bytecode generation needs to happen before Evaluation, because evaluation will return the objects ready to be used.

Evaluation

Transform Exprs into their “usable form”

For example: a new Object, a Var, a Namespace

FnExpr is just `getCompiledClass().newInstance()`

Evaluating means returning the last step of the s-expression transformation process ready to be used in Java code, hence Clojure.

For example, `DefExpr.eval()` returns the created Var after invoking `eval()` on its init instance (for an fn that would be a new instance of an object extending `AFn`).

Compilation

Usually coordination for emit

Compiler.compile namespace -> file

ObjExpr.compile -> coordinate FnExpr emission

Compilation methods tend to coordinate emission action, especially in the case of AOT compilation. Compiler.java contains a compile method that iterates over a list of forms (likely from a clojure script in the classpath) and then analyse, emit and eval each form in them.

Another use of the compile method happens during FnExpr.parser().parse() which is analysing the form. After analysis, FnExpr already contains the generated class.

Emit

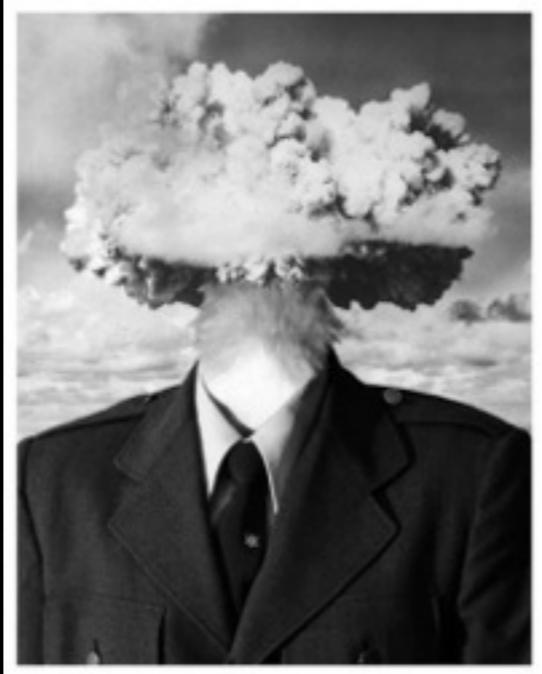
```
(def lister (fn [& args] args))
```

```
DefExpr(  
  Var("lister"),  
  FnExpr(  
    Sym("lister"),  
    PersistentList(  
      FnMethod(  
        LocalBinding(  
          Sym("args")),  
        BodyExpr(  
          PersistentVector(  
            LocalBindingExpr(  
              "args"))))))
```

```
import clojure.lang.RestFn;  
  
public class user$lister extends RestFn {  
  
    public Object doInvoke(Object args) {  
        return args;  
    }  
  
    public int getRequiredArity() {  
        return 0;  
    }  
  
    public user$lister() {  
    }  
}
```

The final step for our lister example consists of invoking eval() on the corresponding DefExpr which essentially is var.bindRoot(init.eval()) where "init" is the init form for this def expression, in our case an FnExpr. init.eval() is compiledClass.newInstance(). At the REPL we can finally invoke the function pointed by the created var.

I know how you feel...



Stress, confusion, hallucinations!

I know how you feel... Let me introduce you to the beasts in more details.

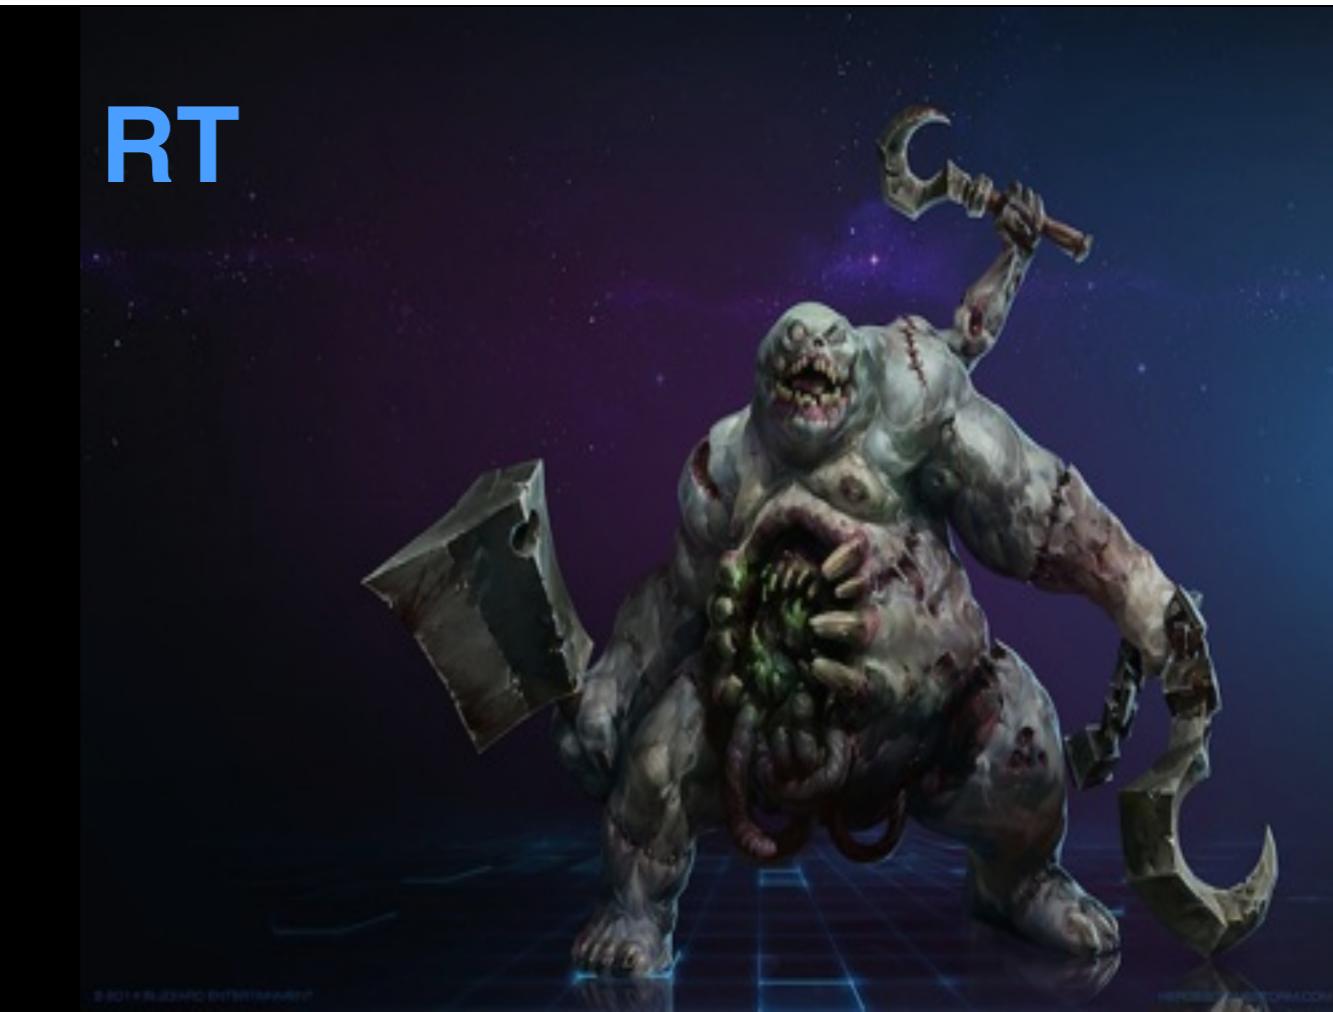
Where are the monsters?

```
package clojure;
import ...
public class main{
    final static private Symbol CLOJURE_MAIN = Symbol.intern("clojure.main");
    final static private Var REQUIRE = RT.var("clojure.core", "require");
    final static private Var LEGACY REPL = RT.var("clojure.main", "legacy-repl");
    final static private Var LEGACY_SCRIPT = RT.var("clojure.main", "legacy-script");
    final static private Var MAIN = RT.var("clojure.main", "main");

    public static void legacy_repl(String[] args) {...}
    public static void legacy_script(String[] args) {...}
    public static void main(String[] args) {
        REQUIRE.invoke(CLOJURE_MAIN);
        MAIN.applyTo(RT.seq(args));
    }
}
```

Our journey begins from the Java main class, the one you invoke to start Clojure with java -jar. This innocent main is hiding the parsing of the entire standard library. As soon as you touch RT its static initialisation kicks in. That in turns trigger other important static inits in Compiler, Var, Namespace.

RT



```
RT.java x

1  /*-
2   * Copyright (c) Rich Mickey. All rights reserved.
3   * The use and distribution terms for this software are covered by the
4   * Eclipse Public License 1.0 (http://opensource.org/licenses/eclipse-1.0.php)
5   * which can be found in the file epl-v1.0.html at the root of this distribution.
6   * By using this software in any fashion, you are agreeing to be bound by
7   * the terms of this license.
8   * You must not remove this notice, or any other, from this software.
9  */
10
11 /* rich Mar 25, 2006 4:28:27 PM */
12
13 package clojure.lang;
14
15 import java.net.MalformedURLException;
16 import java.util.concurrent.atomic.AtomicInteger;
17 import java.util.concurrent.Callable;
18 import java.util.*;
19 import java.util.regex.Matcher;
20 import java.util.regex.Pattern;
21 import java.io.*;
22 import java.lang.reflect.Array;
23 import java.math.BigDecimal;
24 import java.math.BigInteger;
25 import java.security.AccessController;
26 import java.security.PrivilegedAction;
27 import java.net.URL;
28 import java.net.JarURLConnection;
29 import java.nio.charset.Charset;
30
31 public class RT{
32
33     static final public Boolean T = Boolean.TRUE;//Keyword.intern(Symbol.intern(null, "t"));
34     static final public Boolean F = Boolean.FALSE;//Keyword.intern(Symbol.intern(null, "t"));
35     static final public String LOADER_SUFFIX = "__init";
36
37     //simple-symbol->class
38     final static IPersistentMap DEFAULT_IMPORTS = map(
39         //                                         Symbol.intern("RT"), "clojure.lang.RT",
40         //                                         Symbol.intern("Num"), "clojure.lang.Num",
41         //                                         Symbol.intern("Symbol"), "clojure.lang.Symbol",
42         //                                         Symbol.intern("Keyword"), "clojure.lang.Keyword"
43     );

```

RT stands for RunTime, although RT.java today is more an Util class. It's main responsibilities:

- Main Clojure bootstrap
- Basic definitions
- List manipulation

RT <clinit>

```
299 static{
300     Keyword arglistskw = Keyword.intern(null, "arglists");
301     Symbol namesym = Symbol.intern("name");
302     OUT.setTag(Symbol.intern("java.io.Writer"));
303     CURRENT_NS.setTag(Symbol.intern("clojure.lang.Namespace"));
304     AGENT.setMeta(map(DOC_KEY, "The agent currently running an action on this thread, else nil"));
305     AGENT.setTag(Symbol.intern("clojure.lang.Agent"));
306     MATH_CONTEXT.setTag(Symbol.intern("java.math.MathContext"));
307     Var nv = Var.intern(CLOJURE_NS, NAMESPACE, bootNamespace);
308     nv.setMacro();
309     Var v;
310     v = Var.intern(CLOJURE_NS, IN_NAMESPACE, inNamespace);
311     v.setMeta(map(DOC_KEY, "Sets *ns* to the namespace named by the symbol, creating it if needed.",
312                   arglistskw, list(vector(namesym))));
313     v = Var.intern(CLOJURE_NS, LOAD_FILE,
314                     (AFn) invoke(arg1) -> {
315             try {
316                 {
317                     return Compiler.loadFile((String) arg1);
318                 }
319             catch(IOException e)
320             {
321                 throw Util.sneakyThrow(e);
322             }
323         });
324     v.setMeta(map(DOC_KEY, "Sequentially read and evaluate the set of forms contained in the file.",
325                   arglistskw, list(vector(namesym))));
326     try {
327         doInit();
328     }
329     catch(Exception e) {
330         throw Util.sneakyThrow(e);
331     }
332 }
```

There is quite a lot before this static initialisation block, but this is the most important part. Here you can see we are doing the Var.intern of CLOJURE_NS, something we'll be soon using.

But let's move baby steps. The important part here is that doInit() helper. Let's have a deeper look.

RT.dolnit()

```
446 static void doInit() throws ClassNotFoundException, IOException{
447     load("clojure/core");
448
449     Var.pushThreadBindings(
450         RT.mapUniqueKeys(CURRENT_NS, CURRENT_NS.deref(),
451                         WARN_ON_REFLECTION, WARN_ON_REFLECTION.deref(),
452                         RT.UNCHECKED_MATH, RT.UNCHECKED_MATH.deref()));
453     try {
454         Symbol USER = Symbol.intern("user");
455         Symbol CLOJURE = Symbol.intern("clojure.core");
456
457         Var in_ns = var("clojure.core", "in-ns");
458         Var refer = var("clojure.core", "refer");
459         in_ns.invoke(USER);
460         refer.invoke(CLOJURE);
461         maybeLoadResourceScript("user.clj");
462     } finally {
463         Var.popThreadBindings();
464     }
465 }
466 }
```

That's right, we are going to load the entire standard library right there. Also notice the Var.pushThreadBindings(), an idiom that is found often in the Clojure Java sources and clojure standard library as well.

RT.load()

```
414 static public void load(String scriptbase, boolean failIfNotFound) throws IOException, ClassNotFoundException {
415     String classfile = scriptbase + LOADER_SUFFIX + ".class";
416     String cljfile = scriptbase + ".clj";
417     URL classURL = getResource(baseLoader(), classfile);
418     URL cljURL = getResource(baseLoader(), cljfile);
419     boolean loaded = false;
420
421     if((classURL != null &&
422         (cljURL == null
423          || lastModified(classURL, classfile) > lastModified(cljURL, cljfile)))
424        || classURL == null) {
425         try {
426             Var.pushThreadBindings(
427                 RT.mapUniqueKeys(CURRENT_NS, CURRENT_NS.deref(),
428                     WARN_ON_REFLECTION, WARN_ON_REFLECTION.deref(),
429                     RT.UNCHECKED_MATH, RT.UNCHECKED_MATH.deref()));
430             loaded = (loadClassForName(scriptbase.replace('/', '.') + LOADER_SUFFIX) != null);
431         }
432         finally {
433             Var.popThreadBindings();
434         }
435     }
436     if(!loaded && cljURL != null) {
437         if(booleanCast(Compiler.COMPILE_FILES.deref()))
438             compile(cljfile);
439         else
440             loadResourceScript(RT.class, cljfile);
441     }
442     else if(!loaded && failIfNotFound)
443         throw new FileNotFoundException(String.format("Could not locate %s or %s on classpath: ",
444     }
```

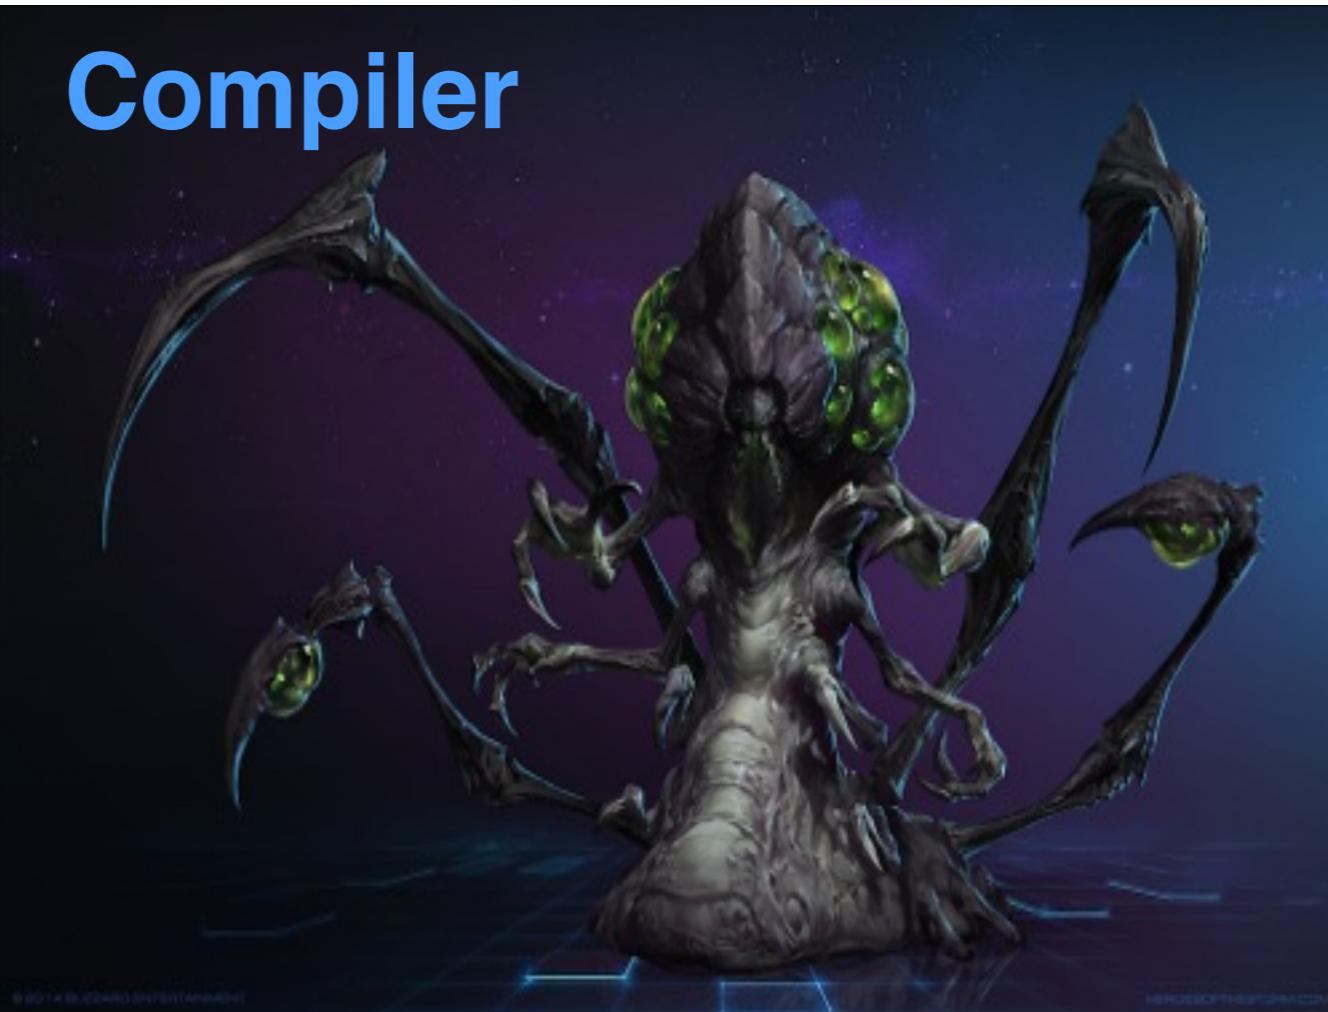
We are not taking the compilation branch here, because we didn't enter from the std lib compile function that is setting that *compile-files* dynamic var to true. Other interesting bits here are searching for an already compiled class for this script base and if existing we are going to skip an additional compilation.

RT.load()

```
364     public static void loadResourceScript(Class c, String name, boolean fail) {
365         int slash = name.lastIndexOf('/');
366         String file = slash >= 0 ? name.substring(slash + 1) : name;
367         InputStream ins = resourceAsStream(baseLoader(), name);
368         if(ins != null) {
369             try {
370                 Compiler.load(new InputStreamReader(ins, UTF8), name, file);
371             } finally {
372                 ins.close();
373             }
374         }
375         else if(failIfNotFound) {
376             throw new FileNotFoundException("Could not locate Clojure resource: " + name);
377         }
378     }
379 }
```

Leaving the RT realm for the moment to enter the compiler.

Compiler



Compiler is a fat class in terms of LOC. The compiler as it is designed today is a “third pass”: first it was conceived as a Lisp emitting Java sources, then translated to Java from Lisp, third it was rewritten to support bytecode emission.

Compiler.java

```
1  /*  
2  * Copyright (c) Rich Hickey. All rights reserved.  
3  * The use and distribution terms for this software are covered by the  
4  * Eclipse Public License 1.0 (http://www.eclipse.org/legal/epl-v10.html).  
5  * By using this software in any fashion, you are agreeing to be bound by  
6  * the terms of this license.  
7  * You must not remove this notice, or any other, from this software.  
8  */  
9  
10 // Rich Aug 21, 2007  
11  
12 package clojure.lang;  
13  
14 //  
15  
16 import clojure.asm.*;  
17 import clojure.asm.commons.GeneratorAdapter;  
18 import clojure.asm.commons.Method;  
19  
20 import java.io.*;  
21 import java.lang.reflect.Constructor;  
22 import java.lang.reflect.Modifier;  
23 import java.util.*;  
24 import java.util.regex.Pattern;  
25 import java.util.regex.Matcher;  
26  
27 //  
28  
29  
30 //import org.objectweb.ESB.*;  
31 //import org.objectweb.ESB.common.Method;  
32 //import org.objectweb.ESB.common.GeneratorAdapter;  
33 //import org.objectweb.ESB.util.TraceClassVisitor;  
34 //import org.objectweb.ESB.util.CheckClassAdapter;  
35 //  
36  
37 public class Compiler implements Opcodes{  
38  
39     static final Symbol DEF = Symbol.intern("def");  
40     static final Symbol LOOP = Symbol.intern("loop");  
41     static final Symbol RECDP = Symbol.intern("recdp");  
42     static final Symbol IF = Symbol.intern("if");  
43     static final Symbol LET = Symbol.intern("let");  
44     static final Symbol LETB = Symbol.intern("letbox");  
45     static final Symbol DO = Symbol.intern("do");  
46     static final Symbol FN = Symbol.intern("fn");  
47     static final Symbol PARCICE = (Symbol) Symbol.intern("//").withMeta(RT.map.Keyword.intern(null, "once"), RT.T);  
48     static final Symbol QUOTE = Symbol.intern("quote");  
49     static final Symbol THE_VAR = Symbol.intern("var");  
50     static final Symbol DOT = Symbol.intern(".");  
51     static final Symbol ASSQIN = Symbol.intern("assq");  
52  
53 //static final Symbol TRY_FINALLY = Symbol.intern("tryFinally")  
}
```

Peculiar style in Clojure sources is the use of inner classes, preferred usually instead of packages. One reason why this can be useful can be found on the design page for JScheme, a 1998 Scheme to Java compiler written by Peter Norvig who clearly stated that in order to make Java coding more "lispy" he hated the idea to always prepend the name of the class to the method invocation. With inner classes, methods also from other classes inside the same compilation unit can be called by name only. Remember there was no import static at that time, same for Rich in 2006.

Compiler.load()

```
7128 public static Object load(Reader rdr, String sourcePath, String sourceName) {
7129     Object EOF = new Object();
7130     Object ret = null;
7131     LineNumberingPushbackReader pushbackReader =
7132         (rdr instanceof LineNumberingPushbackReader) ? (LineNumberingPushbackReader) rdr :
7133         new LineNumberingPushbackReader(rdr);
7134     Var.pushThreadBindings(
7135         RT.mapUniqueKeys(LOADER, RT.makeClassLoader(),
7136             SOURCE_PATH, sourcePath,
7137             SOURCE, sourceName,
7138             METHOD, null,
7139             LOCAL_ENV, null,
7140             LOOP_LOCALS, null,
7141             NEXT_LOCAL_NUM, 0,
7142             RT.READEVAL, RT.T,
7143             RT.CURRENT_NS, RT.CURRENT_NS.deref(),
7144             LINE_BEFORE, pushbackReader.getLineNumber(),
7145             COLUMN_BEFORE, pushbackReader.getColumnNumber(),
7146             LINE_AFTER, pushbackReader.getLineNumber(),
7147             COLUMN_AFTER, pushbackReader.getColumnNumber(),
7148             RT.UNCHECKED_MATH, RT.UNCHECKED_MATH.deref(),
7149             RT.WARN_ON_REFLECTION, RT.WARN_ON_REFLECTION.deref(),
7150             RT.DATA_READERS, RT.DATA_READERS.deref()
7151         ));
7152     try {
7153         for(Object r = LispReader.read(pushbackReader, false, EOF, false); r != EOF;
7154             r = LispReader.read(pushbackReader, false, EOF, false))
7155         {
7156             LINE_AFTER.set(pushbackReader.getLineNumber());
7157             COLUMN_AFTER.set(pushbackReader.getColumnNumber());
7158             ret = eval(r, false);
7159             LINE_BEFORE.set(pushbackReader.getLineNumber());
7160         }
7161     }
```

Load main goal is to loop over the forms found in an input stream, parse them with the reader and pass them down to evaluation, one by one.

Pushing and popping from the thread local storage is not idiomatic Java, where the thread-local is reserved for special usage (like session or transaction IDs). It is pervasive in the Clojure codebase to see values pushed on the thread local and popped right after. Nested invocations will make use of that state.

First we'll introduce another interesting beast, the LispReader.



LispReader

Well, you might ask why there is a LispReader instead of a Clojure reader, aren't we reading Clojure after all? I think it all goes back to when Clojure was implemented as a `clojure.lisp` script of about 1k lines. When the Lisp was ported into Java the name sort of stayed around ([svn://svn.code.sf.net/p/clojure/code/trunk@164](http://svn.code.sf.net/p/clojure/code/trunk@164)).

```
1  /*
2  * Copyright (C) 2004-2009, All Rights Reserved.
3  * The use and distribution terms for this software are covered by the
4  * Eclipse Public License -v 1.0 (http://opensource.org/licenses/eclipse-1.0.php)
5  * which can be found in the file epl-v10.html at the root of this distribution.
6  * By using this software in any fashion, you are agreeing to be bound by
7  * the terms of the EPL.
8  * You must not remove this notice or any other, from this software.
9  */
10
11 package clojure.lang;
12
13 import java.io.EOFException;
14 import java.io.PushbackReader;
15 import java.io.Reader;
16 import java.lang.Character;
17 import java.lang.Class;
18 import java.lang.Exception;
19 import java.lang.IllegalArgumentException;
20 import java.lang.IllegalFormatException;
21 import java.lang.String;
22 import java.lang.StringBuilder;
23 import java.lang.System.currentTimeMillis;
24 import java.lang.Object;
25 import java.lang.RuntimeException;
26 import java.lang.StringBuffer;
27 import java.lang.StringBuilder;
28 import java.lang.Throwable;
29 import java.lang.UnsatisfiedOperationException;
30 import java.util.Collections;
31 import java.util.Map;
32 import java.util.HashMap;
33 import java.util.regex.Matcher;
34 import java.util.regex.Pattern;
35
36 public class LispReader {
37
38     static final Symbol QUOTE = Symbol.intern("QUOTE");
39     static final Symbol DEFN = Symbol.intern("DEFN");
40     static final Symbol DEFNS = Symbol.intern("DEFNS");
41     static final Symbol UNQUOTE = Symbol.intern("UNQUOTE");
42     static final Symbol EVALUTE = Symbol.intern("EVALUTE");
43     static final Symbol EVALUTE_STACKING = Symbol.intern("EVALUTE_STACKING");
44     static final Symbol CONCAT = Symbol.intern("CONCAT");
45     static final Symbol SET = Symbol.intern("SET");
46     static final Symbol LET = Symbol.intern("LET");
47     static final Symbol APP = Symbol.intern("APP");
48     static final Symbol HASHMAP = Symbol.intern("HASHMAP");
49     static final Symbol HASHSET = Symbol.intern("HASHSET");
50     static final Symbol VECTOR = Symbol.intern("VECTOR");
51     static final Symbol MAP = Symbol.intern("MAP");
52
53     static final Symbol METAPRO = Symbol.intern("metapro");
54 }
```

Like the Compiler.java, the LispReader.java contains invokable AFn inner classes, a strategy for each of the thing that can be read. The LispReader is also responsible for a first syntactic check and throws error for malformed maps, vectors, lists and so on.

LispReader.read()

```
151 static public Object read(PushbackReader r, boolean eofIsError, Object eofValue, boolean isRecursive)
152 {
153     if(RT.READEVAL.deref() == UNKNOWN)
154         throw Util.runtimeException("Reading disallowed - *read-eval* bound to :unknown");
155
156     try
157     {
158         for(;;)
159         {
160             int ch = readI(r);
161
162             while(isWhitespace(ch))
163                 ch = readI(r);
164
165             if(ch == -1)
166             {
167                 if(eofIsError)
168                     throw Util.runtimeException("EOF while reading");
169                 return eofValue;
170             }
171
172             if(Character.isDigit(ch))
173             {
174                 Object n = readNumber(r, (char) ch);
175                 if(RT.suppressRead())
176                     return null;
177                 return n;
178             }
179
180             IFn macroFn = getMacro(ch);
181             if(macroFn != null)
182             {
183                 Object ret = macroFn.invoke(r, (char) ch);
```

A special table of reader macros is looked up for the char currently pointed at, if it was not matching other kind of supported literals.

LispReader <clinit>

```
81 static
82 {
83     macros['"'] = new StringReader();
84     macros[';'] = new CommentReader();
85     macros['\''] = new WrappingReader(QOTE);
86     macros['@'] = new WrappingReader(DEREF); //new DerefReader();
87     macros['^'] = new MetaReader();
88     macros['`'] = new SyntaxQuoteReader();
89     macros['~'] = new UnquoteReader();
90     macros['('] = new ListReader();
91     macros[')'] = new UnmatchedDelimiterReader();
92     macros['['] = new VectorReader();
93     macros[']'] = new UnmatchedDelimiterReader();
94     macros['{'] = new MapReader();
95     macros['}'] = new UnmatchedDelimiterReader();
96     macros['|'] = new ArgVectorReader();
97     macros['\\\''] = new CharacterReader();
98     macros['%'] = new ArgReader();
99     macros['#'] = new DispatchReader();
100
101
102     dispatchMacros['^'] = new MetaReader();
103     dispatchMacros['\''] = new VarReader();
104     dispatchMacros['`'] = new RegexReader();
105     dispatchMacros['('] = new FnReader();
106     dispatchMacros['{'] = new SetReader();
107     dispatchMacros['='] = new EvalReader();
108     dispatchMacros['|'] = new CommentReader();
109     dispatchMacros['<'] = new UnreadableReader();
110 }
111 }
```

Using an array of indexes to Reader strategies is a cheap form of dynamic lookup. In order to be extensible without recompilation of the sources, the commonly adopted idiomatic way is a Class.forName of a reader strategy found on the classpath.

LispReader\$FnReader

```
624 public static class FnReader extends AFn{  
625     public Object invoke(Object reader, Object lparen) {  
626         PushbackReader r = (PushbackReader) reader;  
627         if(ARG_ENV.deref() != null)  
628             throw new IllegalStateException("Nested #()s are not allowed");  
629         try  
630         {  
631             Var.pushThreadBindings(  
632                 ARG_ENV, PersistentTreeMap.EMPTY));  
633             unread(r, '(');  
634             Object form = read(r, true, null, true);  
635             PersistentVector args = PersistentVector.EMPTY;  
636             PersistentTreeMap argsyms = (PersistentTreeMap) ARG_ENV.deref();  
637             ISeq rargs = argsyms.rseq();  
638             if(rargs != null)  
639             {  
640                 int higharg = (Integer) ((Map.Entry) rargs.first()).getKey();  
641                 if(higharg > 0)  
642                 {  
643                     for(int i = 1; i <= higharg; ++i)  
644                     {  
645                         Object sym = argsyms.valAt(i);  
646                         if(sym == null)  
647                             sym = garg(i);  
648                         args = args.cons(sym);  
649                     }  
650                 }  
651                 Object restsym = argsyms.valAt(-1);  
652                 if(restsym != null)  
653                 {  
654                     args = args.cons(Compiler._AMP_);  
655                     args = args.cons(restsym);  
656                 }  
657             }  
658         }  
659     }  
660 }
```

Some interesting things in a typical reader strategy are highlighted here. LispReader is passed an instance of a push-back reader, a reader with a cursor that can be moved back and forward. This enables look-ahead of symbols or unread to pass the form to a different reader when necessary. Another important aspect is that beginning from the reader and up to the analyzer, Clojure is eating its own dog food. Persistent data structures are pervasive and Java basic collections are rarely used.

Back to Compiler.load()

```
7128     public static Object load(Reader rdr, String sourcePath, String sourceName) {
7129         Object EOF = new Object();
7130         Object ret = null;
7131         LineNumberingPushbackReader pushbackReader =
7132             (rdr instanceof LineNumberingPushbackReader) ? (LineNumberingPushbackReader) rdr :
7133                 new LineNumberingPushbackReader(rdr);
7134         Var.pushThreadBindings(
7135             RT.mapUniqueKeys(LOADER, RT.makeClassLoader(),
7136                 SOURCE_PATH, sourcePath,
7137                 SOURCE, sourceName,
7138                 METHOD, null,
7139                 LOCAL_ENV, null,
7140                 LOOP_LOCALS, null,
7141                 NEXT_LOCAL_NUM, 0,
7142                 RT.READEVAL, RT.T,
7143                 RT.CURRENT_NS, RT.CURRENT_NS.deref(),
7144                 LINE_BEFORE, pushbackReader.getLineNumber(),
7145                 COLUMN_BEFORE, pushbackReader.getColumnNumber(),
7146                 LINE_AFTER, pushbackReader.getLineNumber(),
7147                 COLUMN_AFTER, pushbackReader.getColumnNumber(),
7148                 RT.UNCHECKED_MATH, RT.UNCHECKED_MATH.deref(),
7149                 RT.WARN_ON_REFLECTION, RT.WARN_ON_REFLECTION.deref(),
7150                 RT.DATA_READERS, RT.DATA_READERS.deref()
7151             ));
7152         try {
7153             for(Object r = LispReader.read(pushbackReader, false, EOF, false); r != EOF;
7154                 r = LispReader.read(pushbackReader, false, EOF, false))
7155             {
7156                 LINE_AFTER.set(pushbackReader.getLineNumber());
7157                 COLUMN_AFTER.set(pushbackReader.getColumnNumber());
7158                 ret = eval(r, false);
7159                 LINE_BEFORE.set(pushbackReader.getLineNumber());
7160             }
7161         }
```

Each read from the LispReader is returning a brand new form as a list of symbols and potentially embedded lists. Despite its controversial name, Compiler.eval() is indeed managing the Analyzer pass.

Compiler.eval()

```
6715
6716
6717
6718
6719
6720
6721
6722
6723
6724
6725
6726
6727
6728
6729
6730
6731
6732
6733
6734
6735
6736
6737
6738
6739
6740
6741
6742
6743
try {
    form = macroexpand(form);
    if(form instanceof ISeq && Util.equals(RT.first(form), DO))
    {
        ISeq s = RT.next(form);
        form = next; i = next; s = RT.next(s)
        eval(RT.first(s), false);
        return eval(next, false);
    }
    else if((form instanceof IType) ||
            (form instanceof IPersistentCollection
             && !(RT.first(form) instanceof Symbol
                   && ((Symbol) RT.first(form)).name.startsWith("def")))) {
        ObjExpr fexpr = (ObjExpr) analyze(C.EXPRESSION, RT.list(FN, PersistentVector.EMPTY, form)
                                         .eval + RT.nextID());
        IFn fn = (IFn) fexpr.eval();
        return fn.invoke();
    }
    else {
        Expr expr = analyze(C.EVAL, form);
        return expr.eval();
    }
} finally
{
    Var.popThreadBindings();
```

We'll recur on Compiler.eval() until we get back something that is not an sequence. As soon as we reach a leaf, we go down analysing it.

analyze()

```
if(form instanceof LazySeq)
{
    form = RT.seq(form);
    if(form == null)
        form = PersistentList.EMPTY;
}
if(form == null)
    return NIL_EXPR;
else if(form == Boolean.TRUE)
    return TRUE_EXPR;
else if(form == Boolean.FALSE)
    return FALSE_EXPR;
Class fclass = form.getClass();
if(fclass == Symbol.class)
    return analyzeSymbol((Symbol) form);
else if(fclass == Keyword.class)
    return registerKeyword((Keyword) form);
else if(form instanceof Number)
    return NumberExpr.parse((Number) form);
else if(fclass == String.class)
    return new StringExpr((String) form).intern();
else if(form instanceof IPersistentCollection && ((IPersistentCollection) form).count() == 0)
{
    Expr ret = new EmptyExpr(form);
    if(RT.meta(form) != null)
        ret = new MetaExpr(ret, MapExpr
            .parse(context == C.EVAL ? context : C.EXPRESSION, ((IObj) form).meta()));
    return ret;
}
else if(form instanceof ISeq)
    return analyzeSeq(context, (ISeq) form, name);
else if(form instanceof IPersistentVector)
    return VectorExpr.parse(context, (IPersistentVector) form);
else if(form instanceof IRecord)
    return new ConstantExpr(form);
else if(form instanceof IType)
    return new ConstantExpr(form);
else if(form instanceof IPersistentMap)
    return MapExpr.parse(context, (IPersistentMap) form);
else if(form instanceof IPersistentSet)
    return SetExpr.parse(context, (IPersistentSet) form);

return new ConstantExpr(form);
```

There are many analysing strategies. One of the most common is analyzeSeq, maybe worth exploring a little further.

```
6649 private static Expr analyzeSeq(C context, ISeq form, String name) {  
6650     Object line = lineDeref();  
6651     Object column = columnDeref();  
6652     if(RT.meta(form) != null && RT.meta(form).containsKey(RT.LINE_KEY))  
6653         line = RT.meta(form).valAt(RT.LINE_KEY);  
6654     if(RT.meta(form) != null && RT.meta(form).containsKey(RT.COLUMN_KEY))  
6655         column = RT.meta(form).valAt(RT.COLUMN_KEY);  
6656     Var.pushThreadBindings(  
6657         RT.map(LINE, line, COLUMN, column));  
6658     try  
6659     {  
6660         Object me = macroexpand1(form);  
6661         if(me != form)  
6662             return analyze(context, me, name);  
6663         Object op = RT.first(form);  
6664         if(op == null)  
6665             throw new IllegalArgumentException("Can't call nil");  
6666         IFn inline = isInline(op, RT.count(RT.next(form)));  
6667         if(inline != null)  
6668             return analyze(context, preserveTag(form, inline.applyTo(RT.n  
6669             IParser p;  
6670             if(op.equals(FN))  
6671                 return FnExpr.parse(context, form, name);  
6672             else if((p = (IParser) specials.valAt(op)) != null)  
6673                 return p.parse(context, form);  
6674             else  
6675                 return InvokeExpr.parse(context, form);  
6676         }  
6677         catch(Throwable e)  
6678         {  
6679             if(!(e instanceof CompilerException))  
6680                 throw new CompilerException((String) SOURCE_PATH.deref(), lin  
6681             else  
6682                 throw (CompilerException) e;  
6683         }  
6684         finally  
6685         {  
6686             Var.popThreadBindings();  
6687         }  
6688     }  
6689 }
```

analyzeSeq()

Sequence analysis performs another round of macro expansion (possibly because it can be invoked from other part of the code) and then figure out what to do based on the first element of the form. If that is fn (quite a common case) we need to parse an FnExpr.

FnExpr parse

```
3902     try
3903     {
3904         fn.compile(n.isVariadic() ? "clojure/lang/RestFn" : "clojure/lang/AFunction",
3905                     (prims.size() == 0)?
3906                         null
3907                         :prims.toArray(new String[prims.size()]),
3908                         fn.onceOnly);
3909     }
3910     catch(IOException e)
3911     {
3912         throw Util.sneakyThrow(e);
3913     }
3914     fn.getCompiledClass();
3915     if(fn.supportsMeta())
3916     {
3917         //System.err.println(name + " supports meta");
3918         return new MetaExpr(fn, MapExpr
3919                             .parse(context == C.EVAL ? context : C.EXPRESSION, fmeta));
3920     }
3921     else
3922     {
3923         return fn;
3924     }
3925 }
```

FnExpr parsing is quite complicated stuff. Something worth noticing is that bytecode generation is triggered here. The call to getCompiledClass() is used for the side effect of having the freshly new generated class loaded in the class loader and ready to use.

```

4093
4094     void compile(String superName, String[] interfaceNames, boolean oneTimeUse) throws IOException{
4095
4096         //...
4097         ClassWriter cw = new ClassWriter(ClassWriter.COMPUTE_MAXS);
4098         ClassWriter cw = new ClassWriter(0);
4099         ClassVisitor cv = cw;
4100
4101         //...
4102         cv.visit(VI_5, ACC_PUBLIC + ACC_SUPER + ACC_FINAL, internalName, null,superName,interfaceNames);
4103
4104         //...
4105         String source = (String) SOURCE.deref();
4106         int lineBefore = (Integer) LINE_BEFORE.deref();
4107         int lineAfter = (Integer) LINE_AFTER.deref() + 1;
4108         int columnBefore = (Integer) COLUMN_BEFORE.deref();
4109         int columnAfter = (Integer) COLUMN_AFTER.deref() + 1;
4110
4111         if(source != null && SOURCE_PATH.deref() != null)
4112         {
4113             //cv.visitSource(source, null);
4114             String smap = "SMAP\n" +
4115                         ((source.lastIndexOf('.') > 0) ?
4116                          source.substring(0, source.lastIndexOf('.'))
4117                          :source)
4118                         //                                         : simpleName)
4119                         + ".java\n" +
4120                         "Clojure\n" +
4121                         "+S Clojure\n" +
4122                         "+F\n" +
4123                         "+ I " + source + "\n" +
4124                         (String) SOURCE_PATH.deref() + "\n" +
4125                         "+L\n" +
4126                         String.format("%d#1,%d:#d\n", lineBefore, lineAfter - lineBefore, lineBefore) +
4127                         "+E";
4128             cv.visitSource(source, smap);
4129         }
4130         addAnnotation(cv, classMeta);
4131         //static fields for constants
4132         for(int i = 0; i < constants.count(); i++)
4133         {
4134             cv.visitField(ACC_PUBLIC + ACC_FINAL
4135                             + ACC_STATIC, constantName(i), constantType(i).getDescriptor(),
4136                             null, null);
4137         }
4138
4139
4140

```

ASM implements a visitor-based generator. Visiting order is important and each visited “context” must be closed when done (the structure of a class can be visualised like visiting a tree where indentation is the depth of the tree). Compilation proceeds top to bottom, emitting static context first (declaration, initialisation), constructors and then methods.

Bytecode generation needs to iterate over the different arities of the Clojure function and their closed overs, generating static variables and instance variables to hold references to other functions.

FnExpr.emitMethods()

```
3766 protected void emitMethods(ClassVisitor cv){  
3767     //override of invoke/doInvoke for each method  
3768     for(ISeq s = RT.seq(methods); s != null; s = s.next())  
3769     {  
3770         ObjMethod method = (ObjMethod) s.first();  
3771         method.emit(this, cv);  
3772     }  
3773  
3774     if(isVariadic())  
3775     {  
3776         GeneratorAdapter gen = new GeneratorAdapter(ACC_PUBLIC,  
3777                                         Method.getMethod("int getRequiredArity()")  
3778                                         ,  
3779                                         null,  
3780                                         null,  
3781                                         cv);  
3782         gen.visitCode();  
3783         gen.push(variadicMethod.reqParms.count());  
3784         gen.returnValue();  
3785         gen.endMethod();  
3786     }  
3787 }
```

There is a lot more happening on `ObjExpression.compile()` before each arity method can be emitted. Static initialisation, constructors and initialisation of instance state.

`emitMethods()` manages emission of bytecode for each arity of the Clojure function. It just delegates to each `ObjMethod` instance and emits also the `getRequiredArity()` helper to return how many parameters are requested.

ObjMethod.emit() FnMethod.doEmit()

```
5589 public void emit(ObjExpr fn, ClassVisitor cv){  
5590     Method m = new Method(getMethodName(), getReturnType(), getArgTypes());  
5591  
5592     GeneratorAdapter gen = new GeneratorAdapter(ACC_PUBLIC,  
5593             m,  
5594             null,  
5595             //todo don't hardwire this  
5596             EXCEPTION_TYPES,  
5597             cv);  
5598     gen.visitCode();  
5599  
5600     Label loopLabel = gen.mark();  
5601     gen.visitLineNumber(line, loopLabel);  
5602     try  
5603     {  
5604         Var.pushThreadBindings(RT.map(LOOP_LABEL, loopLabel, METHOD, this));  
5605         body.emit(C.RETURN, fn, gen);  
5606         Label end = gen.mark();  
5607         gen.visitLocalVariable("this", "Ljava/lang/Object;", null, loopLabel, end, 0);  
5608         for(ISeq lbs = argLocals.seq(); lbs != null; lbs = lbs.next())  
5609         {  
5610             LocalBinding lb = (LocalBinding) lbs.first();  
5611             gen.visitLocalVariable(lb.name, "Ljava/lang/Object;", null, loopLabel, end,  
5612             );  
5613         }  
5614     }  
5615     finally  
5616     {  
5617         Var.popThreadBindings();  
5618     }
```

Each ObjMethod contains a body instance that is requested to emit in turn. BodyExpr is basically a wrapper for an implicit do form which contains any number of instructions. Around the instructions ObjMethod takes care of assigning local variables.

BodyExpr.emit()

The screenshot shows a Java code editor with the following code snippet:

```
5849     public void emit(C context, ObjExpr objx, GeneratorAdapter gen){  
5850         for(int i = 0; i < exprs.count() - 1; i++)  
5851         {  
5852             Expr e = (Expr) exprs.nth(i);  
5853             e.emit(C.STATEMENT, objx, gen);  
5854         }  
5855         Expr last = exprs.last();  
5856         last.emit(C.STATEMENT, objx, gen);  
5857     }  
5858     public boolean hasJavaClass()  
5859     {  
5860         return true;  
5861     }  
5862     public Class<?> getJavaClass()  
5863     {  
5864         return null;  
5865     }  
5866     private Expr last;  
5867 }  
5868 public static class BodyExpr extends Expr  
5869 {  
5870     public static class BodyExprBuilder  
5871     {  
5872         public static BodyExpr build(BodyExprBuilder builder)  
5873         {  
5874             return null;  
5875         }  
5876     }  
5877 }  
5878 }  
5879 }  
5880 }  
5881 }  
5882 }  
5883 }  
5884 }  
5885 }  
5886 }  
5887 }  
5888 }  
5889 }  
5890 }  
5891 }  
5892 }  
5893 }  
5894 }  
5895 }  
5896 }  
5897 }  
5898 }  
5899 }  
5900 }  
5901 }  
5902 }  
5903 }  
5904 }
```

A tooltip window titled "Choose Implementation of emit (40 found)" is displayed over the line of code "e.emit(C.STATEMENT, objx, gen);". The tooltip lists 40 different Expr classes from the clojure.lang package, each with a "closure" icon and a file path. The "BodyExpr" entry is highlighted.

BodyExpr.emit() iterates over the collected expressions in turn (the body of a method being an implicit “do form” if none specified).

The list of potential implementors any of the Exprs classes (54).

The bottom! LocalBindingExpr.emit()

```
5757 public boolean canEmitPrimitive(){  
5758     return b.getPrimitiveType() != null;  
5759 }  
5760  
5761 public void emitUnboxed(C context, ObjExpr objx, GeneratorAdapter gen){  
5762     objx.emitUnboxedLocal(gen, b);  
5763 }  
5764  
5765 public void emit(C context, ObjExpr objx, GeneratorAdapter gen){  
5766     if(context != C.STATEMENT)  
5767         objx.emitLocal(gen, b, shouldClear);  
5768 }  
5769  
5770 public Object evalAssign(Expr val) throws new UnsupportedOperationException("Can't  
5771  
5772 public void emitAssign(C context, ObjExpr objx, GeneratorAdapter gen, Expr val){  
5773     objx.emitAssignLocal(gen, b, val);  
5774     if(context != C.STATEMENT)  
5775         objx.emitLocal(gen, b, false);  
5776 }  
5777 }  
5778 }
```

Finally we reached a leaf! the place where args (a local close over that was collected from the LispReader as a list of arguments) is taken and emitted. It is not actually returned here, it needs another specific ASM instruction to close the method.

Quick demo

We're just going to show everything together, evaluating the "lister" function and dumping the generated bytecode to the file system, so we can decompile it and have a look inside.

DynamicClassLoader



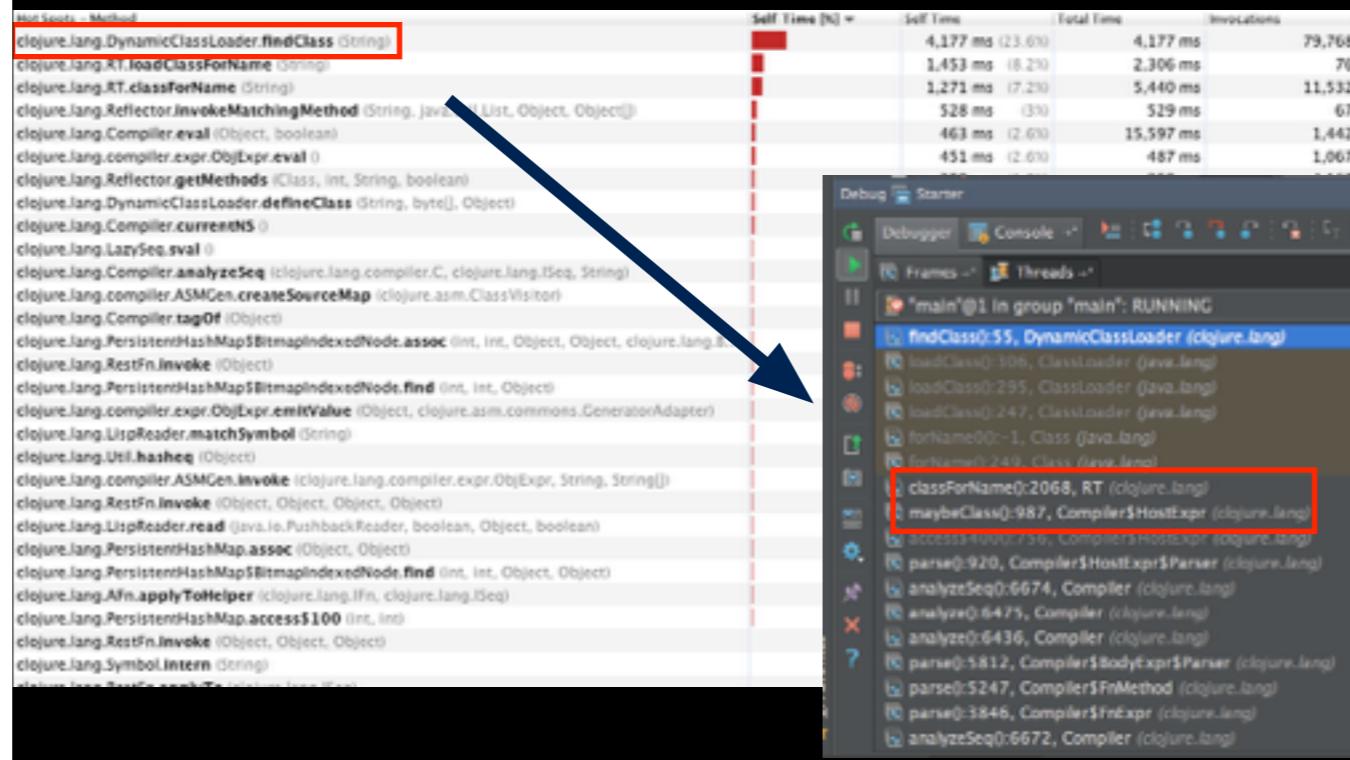
Stressing the compiler

Hot Spots - Method	Self Time [ms]	Self Time	Total Time	Invocations
clojure.lang.Compiler\$HostExpr.maybeClass (Object, boolean)	13,308 ms	(3.1%)	13,349 ms	39,189
clojure.lang.Compiler.eval (Object, boolean)	2,648 ms	(8.6%)	30,799 ms	5,418
clojure.lang.Compiler\$ObjExpr.eval ()	1,980 ms	(6.4%)	1,981 ms	4,200
clojure.lang.Compiler.macropexpand1 (Object)	1,470 ms	(4.8%)	6,874 ms	108,806
clojure.lang.Compiler\$ObjExpr.getCompiledClass ()	1,313 ms	(4.3%)	1,313 ms	10,979
clojure.lang.Compiler.load (java.io.Reader, String, String)	923 ms	(3.1%)	30,845 ms	214
clojure.lang.Compiler.currentNS ()	795 ms	(2.6%)	795 ms	296,779
clojure.lang.Compiler\$ObjExpr.compile (String, String[], boolean)	771 ms	(2.5%)	3,620 ms	6,779
clojure.lang.Compiler.analyzeSeq (clojure.lang.Compiler.C, clojure.lang.ISeq, String)	736 ms	(2.4%)	23,952 ms	98,397
clojure.lang.Compiler.tagOf (Object)	547 ms	(1.8%)	547 ms	236,593
clojure.lang.Compiler.analyze (clojure.lang.Compiler.C, Object, String)	347 ms	(1.1%)	23,959 ms	210,337
clojure.lang.Compiler\$FnExpr.parse (clojure.lang.Compiler.C, clojure.lang.ISeq, String)	310 ms	(1.0%)	23,425 ms	6,616
clojure.lang.Compiler\$ObjExpr.emitValue (Object, clojure.asm.commons.GeneratorAdapter)	278 ms	(0.9%)	307 ms	55,315
clojure.lang.Compiler.lookupVar (clojure.lang.Symbol, boolean, boolean)	222 ms	(0.7%)	879 ms	133,443
clojure.lang.Compiler\$instanceMethodExpr.eval ()	221 ms	(0.7%)	221 ms	2
clojure.lang.Compiler\$InvokeExpr.eval ()	193 ms	(0.6%)	373 ms	151
clojure.lang.Compiler\$HostExpr\$Parser.parse (clojure.lang.Compiler.C, Object)	173 ms	(0.6%)	7,966 ms	10,206
clojure.lang.Compiler\$StaticMethodExpr.<init> (String, Int, Int, clojure.lang.Symbol, Class, String,...)	158 ms	(0.5%)	605 ms	5,981
clojure.lang.Compiler.resolveIn (clojure.lang.Namespace, clojure.lang.Symbol, boolean)	157 ms	(0.5%)	214 ms	28,454
clojure.lang.Compiler\$StaticMethodExpr.eval ()	155 ms	(0.5%)	155 ms	442
clojure.lang.Compiler.analyzeSymbol (clojure.lang.Symbol)	140 ms	(0.5%)	1,860 ms	63,312
clojure.lang.Compiler\$InvokeExpr.parse (clojure.lang.Compiler.C, clojure.lang.ISeq)	136 ms	(0.4%)	11,203 ms	27,132
clojure.lang.Compiler\$instanceMethodExpr.<init> (String, Int, Int, clojure.lang.Symbol, clojure.lang...)	128 ms	(0.4%)	206 ms	3,732
clojure.lang.Compiler\$ObjExpr.emitConstants (clojure.asm.commons.GeneratorAdapter)	123 ms	(0.4%)	477 ms	5,614
clojure.lang.Compiler.registerLocal (clojure.lang.Symbol, clojure.lang.Symbol, clojure.lang.Compl...)	123 ms	(0.4%)	644 ms	20,636
clojure.lang.Compiler\$FnMethod.parse (clojure.lang.Compiler.ObjExpr, clojure.lang.ISeq, boolean)	120 ms	(0.4%)	20,855 ms	6,923
clojure.lang.Compiler.nameSpaceFor (clojure.lang.Namespace, clojure.lang.Symbol)	105 ms	(0.3%)	105 ms	116,151
clojure.lang.Compiler\$ObjExpr.constantType (int)	104 ms	(0.3%)	104 ms	142,887
Total time: 13,349 ms (3.1%)				

Stressing the compiler isn't that difficult. Loading core.clj is probably all that it takes. Just in case, feed it a project with many dependencies and watch the compiler at work while AOT compiling the entire thing.

Some low hanging fruits here (as of 1.6): why that HostExpr.maybeClass() is taking way more than anything else? Little caching might be helpful.

clojure.lang.* at work



Same profiling, but the entire clojure.lang Java package. The DynamicClassLoader.findClass() is quite stressed during compilation. It turns out it is again HostExpr.maybeClass() invoking a RT.className, then Class.forName that triggers the findClass on the custom classloader.

HostExpr is just one example of code that needs to lookup classes, nothing wrong with that. It just happens that is showing that searching and defining classes is a custom feature handled by the DynamicClassLoader.

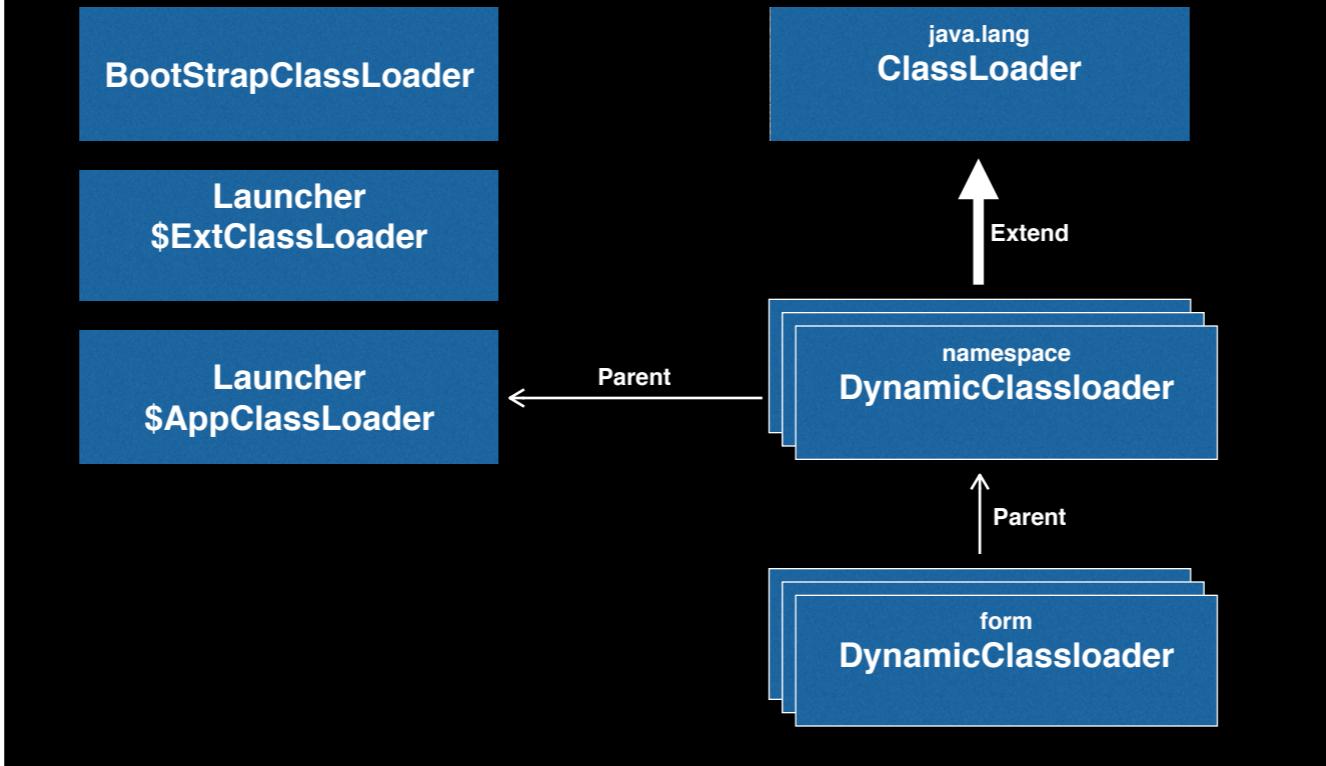
HostExpr.maybeClass()

```
966     private static Class maybeClass(Object form, boolean stringOk) {
967         if(form instanceof Class)
968             return (Class) form;
969         Class c = null;
970         if(form instanceof Symbol)
971         {
972             Symbol sym = (Symbol) form;
973             if(sym.ns == null) //if ns-qualified can't be classname
974             {
975                 if(Util.equals(sym, COMPILE_STUB_SYM.get()))
976                     return (Class) COMPILE_STUB_CLASS.get();
977                 if(svm.name.indexOf('.') > 0 || svm.name.charAt(0) == '[')
978                     c = RT.className(sym.name);
979                 else
980                 {
981                     Object o = currentNS().getMapping(sym);
982                     if(o instanceof Class)
983                         c = (Class) o;
984                     else
985                     {
986                         try{
987                             c = RT.className(sym.name);
988                         }
989                         catch(Exception e){
990                             //...
991                         }
992                     }
993                 }
994             }
995         }
996     }
997     else if(stringOk && form instanceof String)
998         c = RT.className((String) form);
999
1000 }
```

Java interop is used extensively during Clojure bootstrapping. HostExpr parser is delegated to deal with the parsing of the Class or instance target for a method invocation. Maybe class is used to determine if the Symbol is indeed a class and transform it into a real one.

RT.className is invoking Class.forName() inside. This in turn mutates into a findClass() call for all the class loaders in the hierarchy. Consider the hundreds of them instantiated, the Compiler is spending quite an amount of time searching for them.

Pushing classloaders



So, what's going on here? A little bit of Clojure class loaders organisation.

Every time a form is evaluated, at the REPL or during AOT compilation, a new `DynamicClassLoader` instance is created to hold the generated classes related to the bytecode compilation of that form. This form-classloader is child of another `DynamicClassLoader` instance that was created for the namespace the form is defined into. The namespace-classloader is then in turn child of the standard application classloader.

Once a namespace is required as the natural chain of dependencies during compilation it will be subject to the process of class generation for each function contained in them, **doesn't matter if they are used or not**.



This is what happens if I stick a couple of print lines to show every time a class loader is created and all the time a class loader instance is searched for a class implementation.

Luckily for us, this is only a bootstrap problem. None of that happens for our AOT compiled apps.

Supporting Dynamicity

```
public Class defineClass(String name, byte[] bytes, Object srcForm){  
    Util.clearCache(rq, classCache);  
    Class c = defineClass(name, bytes, 0, bytes.length);  
    classCache.put(name, new SoftReference(c,rq));  
    return c;  
}  
  
protected Class<?> findClass(String name) throws ClassNotFoundException{  
    Reference<Class> cr = classCache.get(name);  
    if(cr != null)  
    {  
        Class c = cr.get();  
        if(c != null)  
            return c;  
        else  
            classCache.remove(name, cr);  
    }  
    return super.findClass(name);  
}
```

Clojure classloading is there to support the dynamicity of Clojure, the fact that functions or entire namespaces can be thrown away and redefined at anytime. The class loader that was generated for a namespace and all its form-classloaders is stored in a Var and when a reloaded is required all it takes to remove their only reference is just an “alter-var-root” away.

When the var root holding a reference to the class loader is gone, the class loader can be garbage collected and along with that all the references to classes that are unused.

Since some of the function generated during the namespace compilation might never be used, it makes sense to give the garbage collector a chance to collect them. That's why they are wrapped in SoftReferences. If PermGen is close to be exhausted, the JVM can claim soft references and free up some space.

Resources

All snippets are from this SHA of [clojure fastload branch](#)

The [Clojure mailing list](#) especially older posts

The [Clojure IRC channel logs](#)

Interesting [interview](#) with Rich about Clojure history

An [old](#) but still useful class diagram

Very nice [blog coverage](#) about Lisp50

Presenter notes on these slides!

I so miss the time when Rich was discussing about everything Clojure on the #clojure IRC channel!

Bonus: Clojure in Lisp

```
1 ;/**/  
2 ; * Copyright (c) Rich Hickey. All rights reserved.  
3 ; * The use and distribution terms for this software are covered by the  
4 ; * Common Public License 1.0 (http://opensource.org/licenses/cpl.php)  
5 ; * which may be found in the file CPL-V1.0.txt at the root of this distribution.  
6 ; * By using this software in any fashion, you are agreeing to be bound by  
7 ; * the terms of this license.  
8 ; * You must not remove this notice, or any other, from this software.  
9 ; **/  
10 (defpackage "clojure"  
11   (:export :load-types :*namespace-separator*  
12             :newobj :@ :compile-to :*clojure-source-path* :*clojure-target-path*  
13             :in-module"  
14             "defn" "def" "fn"  
15             "if" "and" "or" "not" "when" "unless"  
16             "block" "let" "let*" "letfn"  
17             "set" "pset" "set*" "do"  
18             "try" "ex"  
19             "char" "boolean" "byte" "short" "int" "long" "float" "double"  
20             "import"))  
21  
22 (in-package "clojure")  
23  
24 (defvar *namespace-separator* nil  
25   "set to #\: for JVM, #. for CLI")  
26  
27  
28 (defconstant +MAX-POSITIONAL-ARITY+ 5)  
29  
30 (defvar *host* nil) ; :jvm or :cli  
31 (defvar *clojure-source-path*)  
32 (defvar *clojure-target-path*)  
33 (defvar *symbols*)  
34 (defvar *keywords*)  
35 (defvar *vars*)  
36 (defvar *accessors*)  
37 (defvar *defvars*)  
38 (defvar *defns*)  
39 (defvar *quoted-aggregates* nil)  
40 (defvar *nested-fn-bindings*)  
41 (defvar *var-env* nil)  
42 (defvar *frame* nil)  
43 (defvar *next-id*)  
44  
45 (defvar *imports*)  
46  
47
```

Bonus: Lisp2Java

```
/** Copyright (c) Rich Hickey. All rights reserved.  
 * The use and distribution terms for this software are covered by the  
 * Eclipse Public License 1.0 (http://opensource.org/licenses/eclipse-1.0.php)  
 * which can be found in the file LICENSE at the root of this distribution.  
 * By using this software or library you are agreeing to the terms  
 * the terms of the Eclipse Public License.  
 * You must not remove this notice, or any other, from this software.  
 **/  
  
/* rich Aug 21, 2007 */  
  
package clojure.lang;  
  
/*  
  
import clojure.asm.*;  
import clojure.asm.commons.GeneratorAdapter;  
import clojure.asm.commons.Method;  
  
import java.io.*;  
import java.lang.reflect.Constructor;  
import java.lang.reflect.Modifier;  
import java.util.*;  
import java.util.regex.Pattern;  
import java.util.regex.Matcher;  
  
/**/  
  
//import org.objectweb.asm.*;  
//import org.objectweb.asm.commons.Method;  
//import org.objectweb.asm.commons.GeneratorAdapter;  
//import org.objectweb.asm.util.TraceClassVisitor;  
//import org.objectweb.asm.util.CheckClassAdapter;  
//  
public class Compiler implements Opcodes{  
  
static final Symbol DEF = Symbol.intern("def");  
static final Symbol LOOP = Symbol.intern("loop*");  
static final Symbol RECUR = Symbol.intern("recur");  
static final Symbol IF = Symbol.intern("if");  
static final Symbol LET = Symbol.intern("let*");  
static final Symbol LETFN = Symbol.intern("letfn*");  
static final Symbol DO = Symbol.intern("do");
```



```
/* Generated by Clojure */
namespace clojure.lib {
using System;
using org.clojure.runtime;

public class Clojure{
/* Generated by Clojure from the following Lisp:
(defn* apply ((fn & args+)) (.applyTo fn __tld (spread* args+)))
*/
    public class apply : RestFn1{
override public Object doInvoke(ThreadLocalData __tld, Object fn, Cons argsPLUS__)
{
if(__tld == null) __tld = ThreadLocalData.get();
return ((IFn)ACC_clojure_applyTo).invoke(__tld, fn, clojure__tld.getValue(__tld), clojure_spreadSTAR);
}
}
/* Generated by Clojure from the following Lisp:
(defn* complement ((fn) (fn (& args) (not (apply fn args)))))
*/
    public class complement : AFn{
override public Object invoke(ThreadLocalData __tld, Object fn)
{
return (new FN_2163(fn));
}
    public class FN_2163 : RestFn0{
Object fn;
public FN_2163 (Object fn){
this.fn = fn;
}
override public Object doInvoke(ThreadLocalData __tld, Cons args)
{
return ((clojure_apply.fn.invoke(__tld, fn, args))==null?RT.T:null);
}
}
}
/* Generated by Clojure from the following Lisp:
(defn* constantly ((x) (fn (& args) x)))

```

Bonus: Lisp2C#

Special thanks

The entire “Clojure Table” at the **MailOnline**

Phil Potter and the other EuroClojure speakers

<http://mr--jack.deviantart.com> incredible monsters

and finally Rich for the great learning opportunity

By the way, **WE ARE HIRING!**